

Diversity and Inclusion

We strive to host inclusive, accessible events that enable all individuals to engage fully.

If you have a concern, if you need an accommodation for a disability, or if you would like to learn more, please talk to one of our Microsoft staff on-site and we'll do our best to assist you.

Thank you.

Practice
Awareness



Exercise
Curiosity



Demonstrate
Courage






GPU Accelerated IoT Workloads at the Edge

Paul DeCarlo
Principal Cloud Advocate
@pjdecarlo

Presentation available @
aka.ms/intelligentedgeholdeck



Demo

YoloV3 IoT Edge Module Running on Jetson Xavier

YoloV3 IoT Edge Module Running on Jetson Xavier



Agenda



Introduce Azure IoT Edge



Introduce Nvidia Jetson Devices




Developing IoT Edge Modules with GPU acceleration



Intelligent Edge Hands-On-Lab




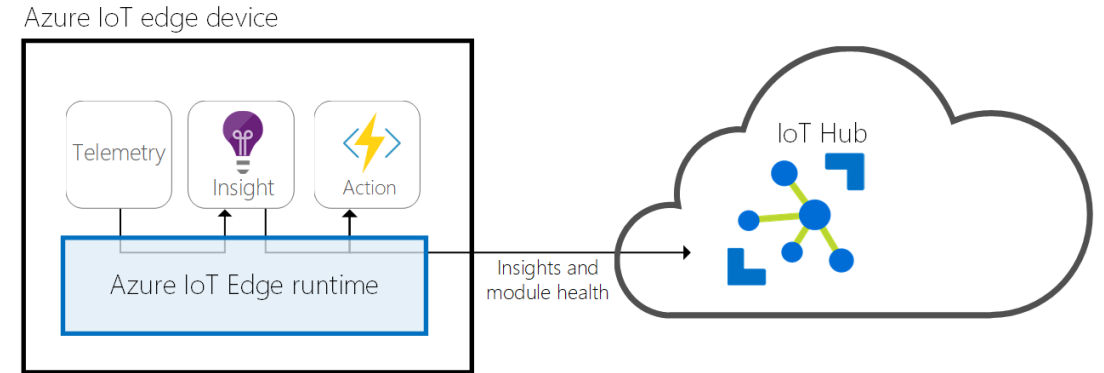
Additional Resources



Azure IoT Edge

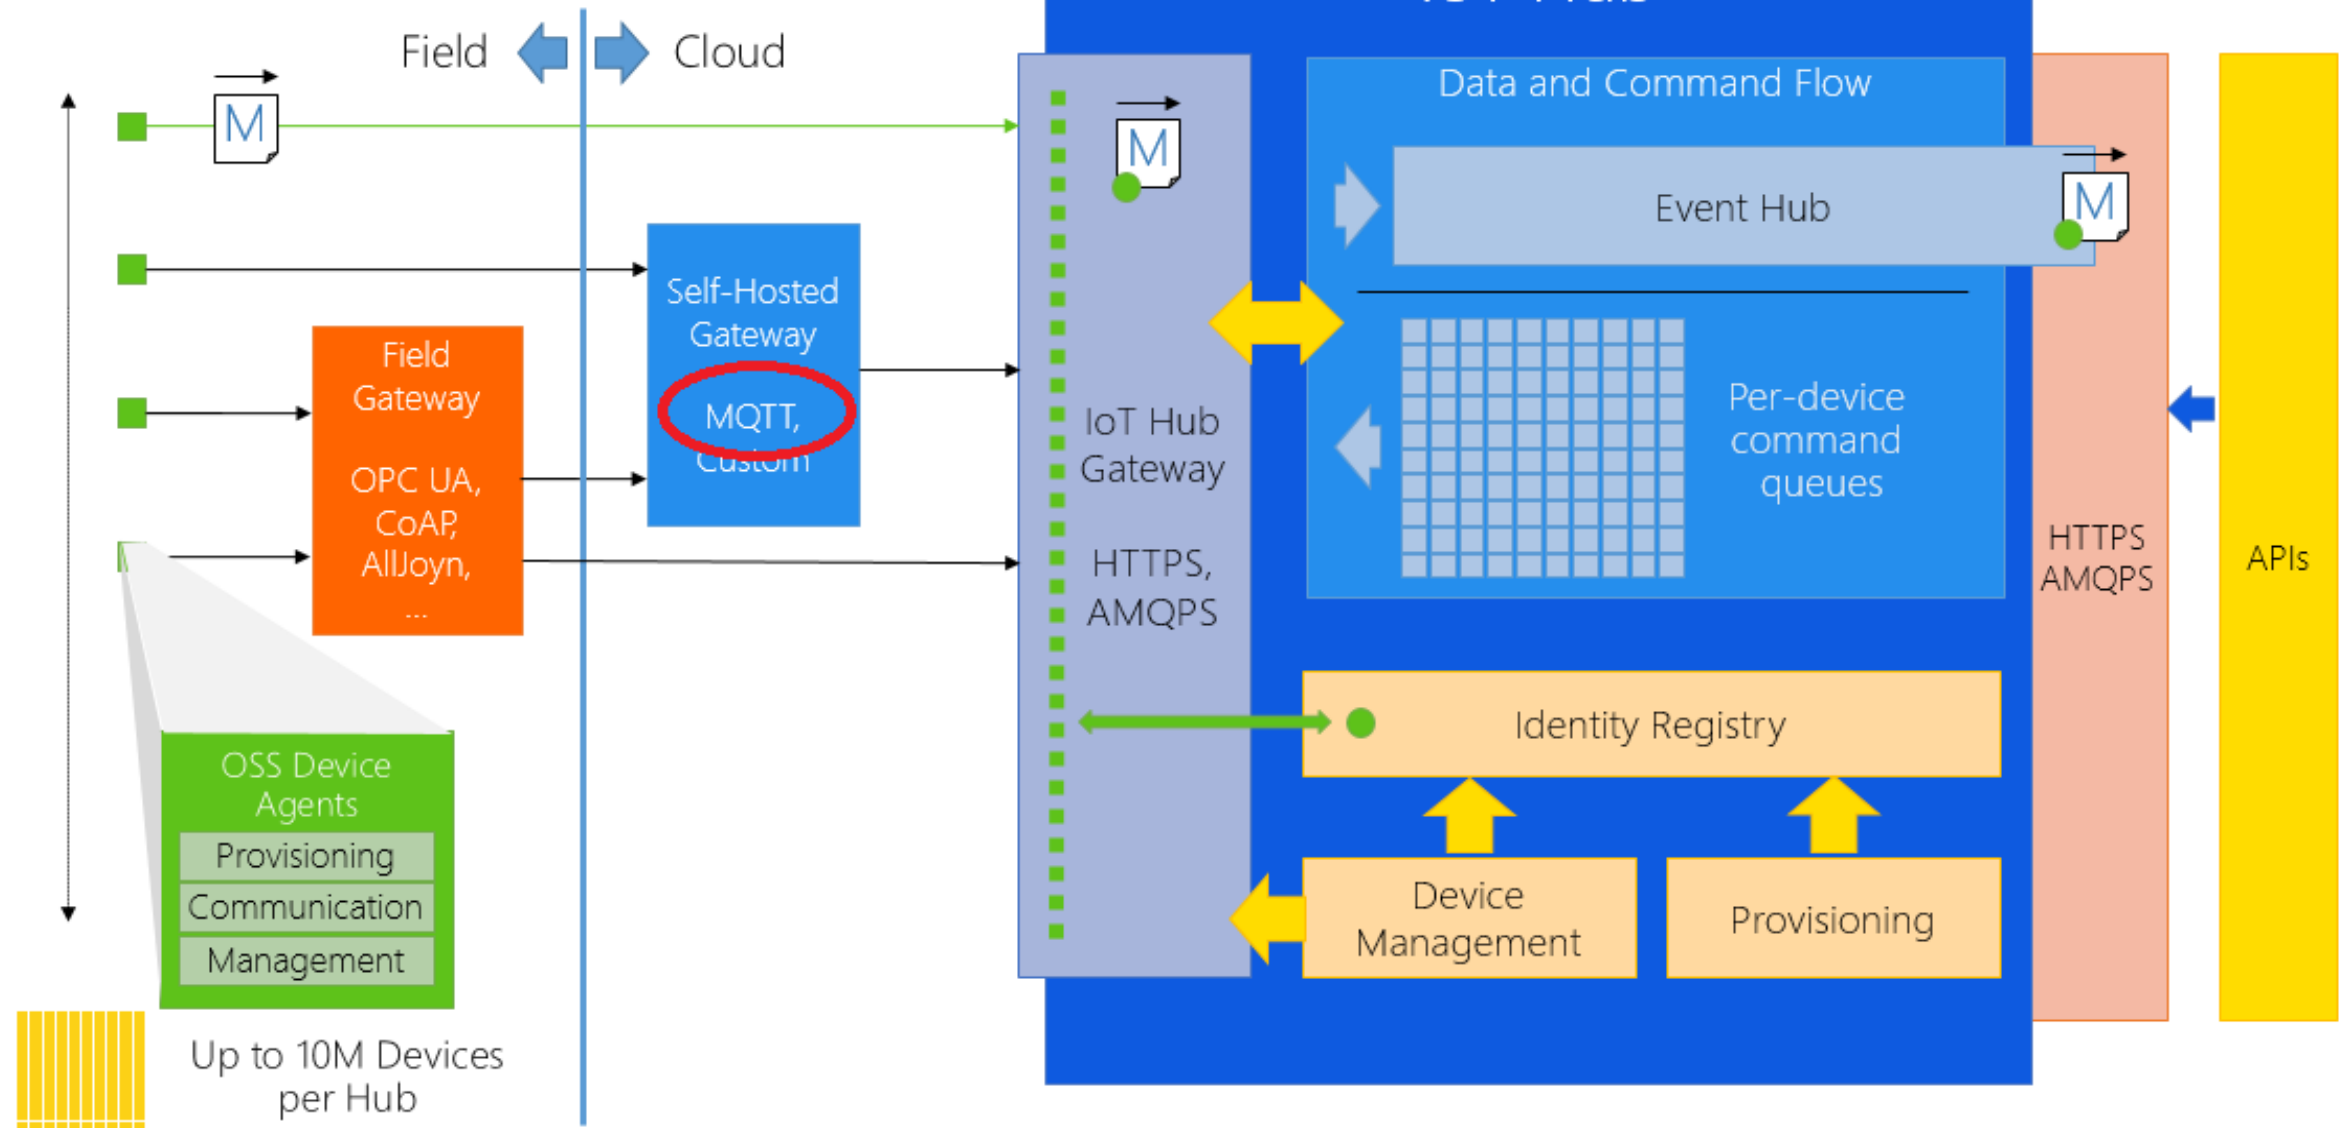
- Built on container technology as 'modules'
- Modules support Python, NodeJS, .Net Core, Java, & C
- Low-latency AMQP / MQTT data transport
- Operate in offline / intermittent network conditions
- Supports Linux X64 | ARM32/64 , Windows X64

OSS and available @ <https://github.com/Azure/iotedge>



Azure IoT Hub


Field ↔ Cloud



System Modules

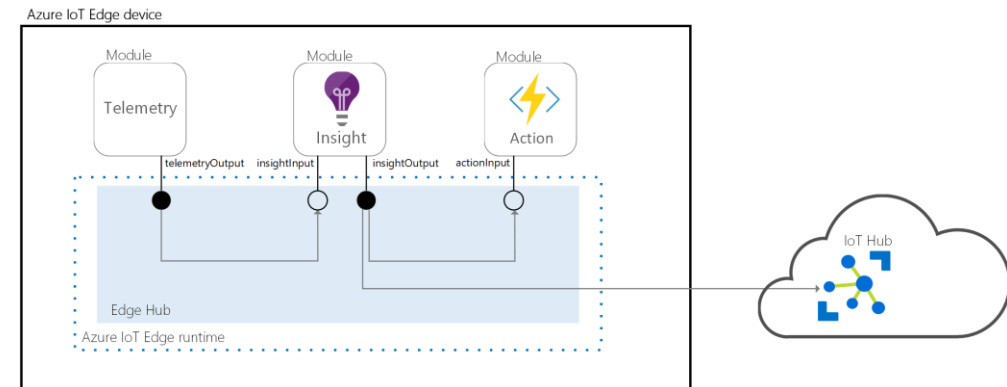
edge-agent:

Deployment & Container orchestration
Ensures module uptime



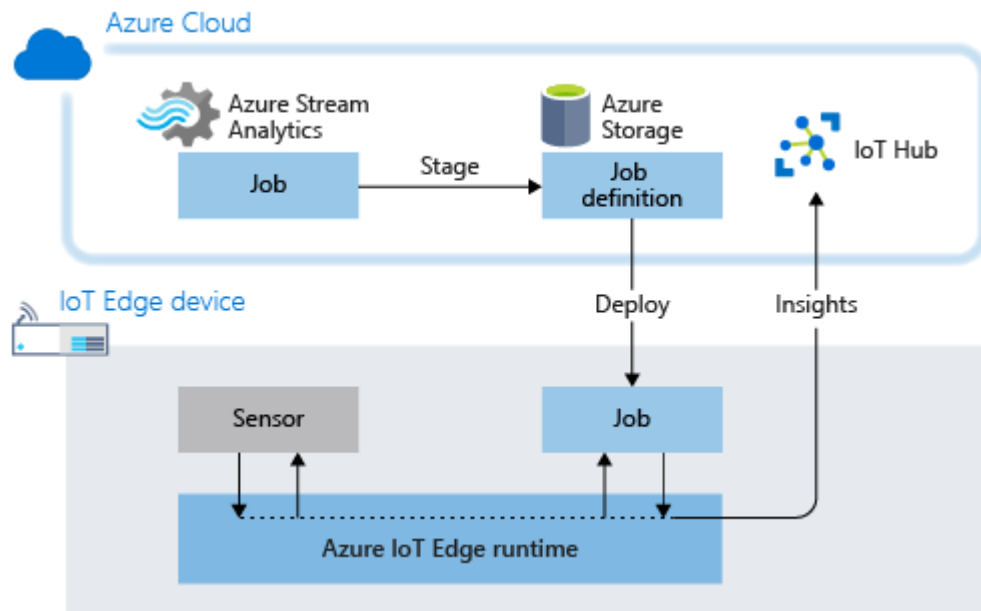
edge-hub:

Communication to/from Azure IoT Hub
Inter-module communication

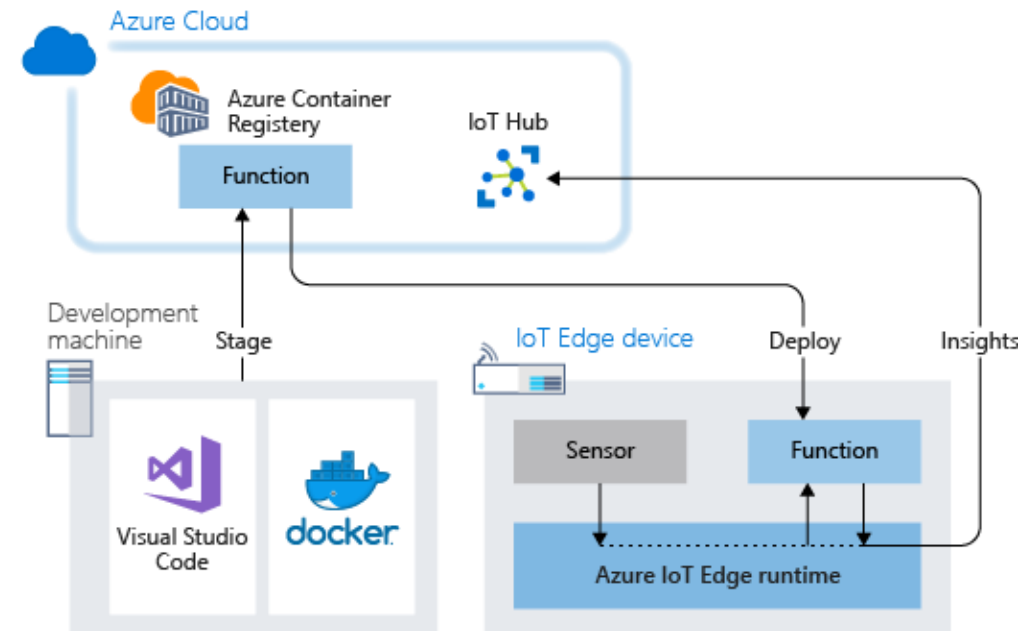


Azure Services as Edge Modules

Azure Stream Analytics

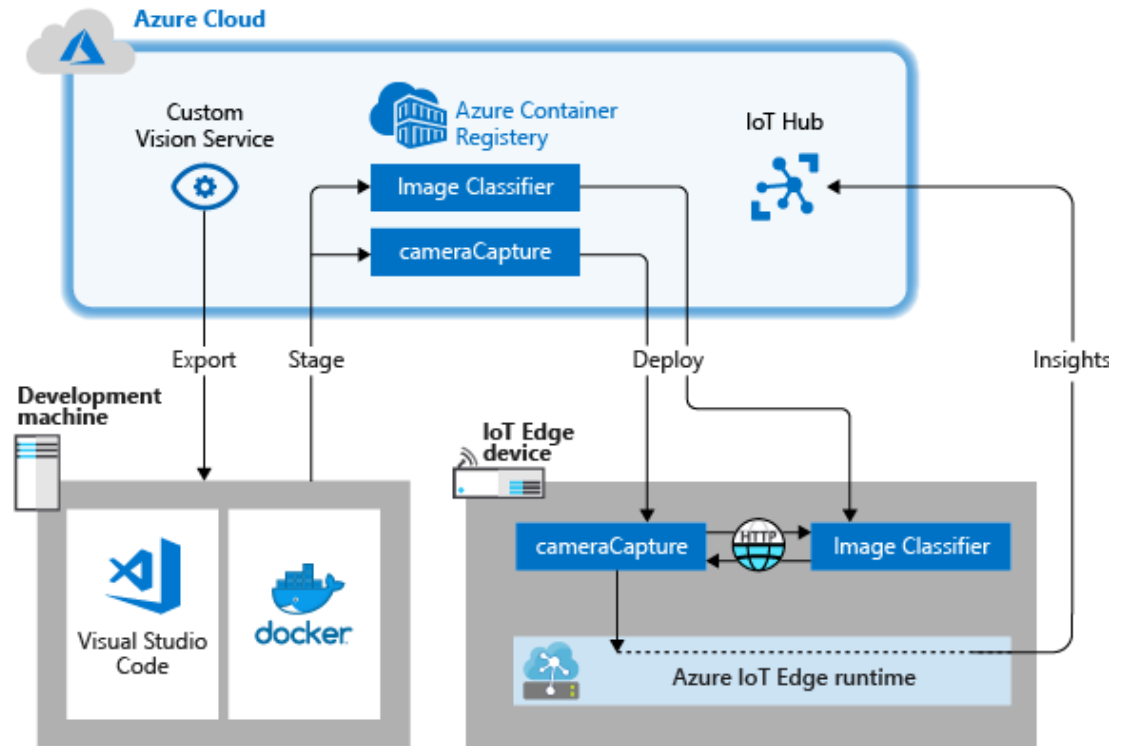


Azure Functions



AI/ML Modules



- Cognitive Services including:
 - Text Analytics - Key Phrase Extraction, Language Detection, Sentiment Analysis
 - Face – Detection, Verification, and Emotion
 - Computer Vision – Text Recognition, custom vision module allows to finds whatever you train it to find
- Azure Machine Learning at the Edge
 - Export AML models to run a Edge modules





IoT Edge Deployment Process



Publish Deployment to IoT Hub for targeted devices



IoT Edge Deployment created



IoT Edge Devices matching target criteria in respective DeviceTwin receive updated deployment configuration



IoT Edge Deployments

An IoT Edge deployment configures a target set of IoT Edge devices to run a set of IoT Edge modules. Each deployment continuously monitors the target condition it meets. Each IoT Hub supports up to 10 deployments. [Learn more about IoT Edge deployment.](#)

ID	TARGET CONDITION
vsts-deployment-cpu	tags.environment='dev' and tags.platform='cpu'

Target Condition

The target condition defines the devices to be targeted.

* Target Condition ⓘ

tags.environment='dev' and tags.platform='cpu'

Show devices Reset target condition

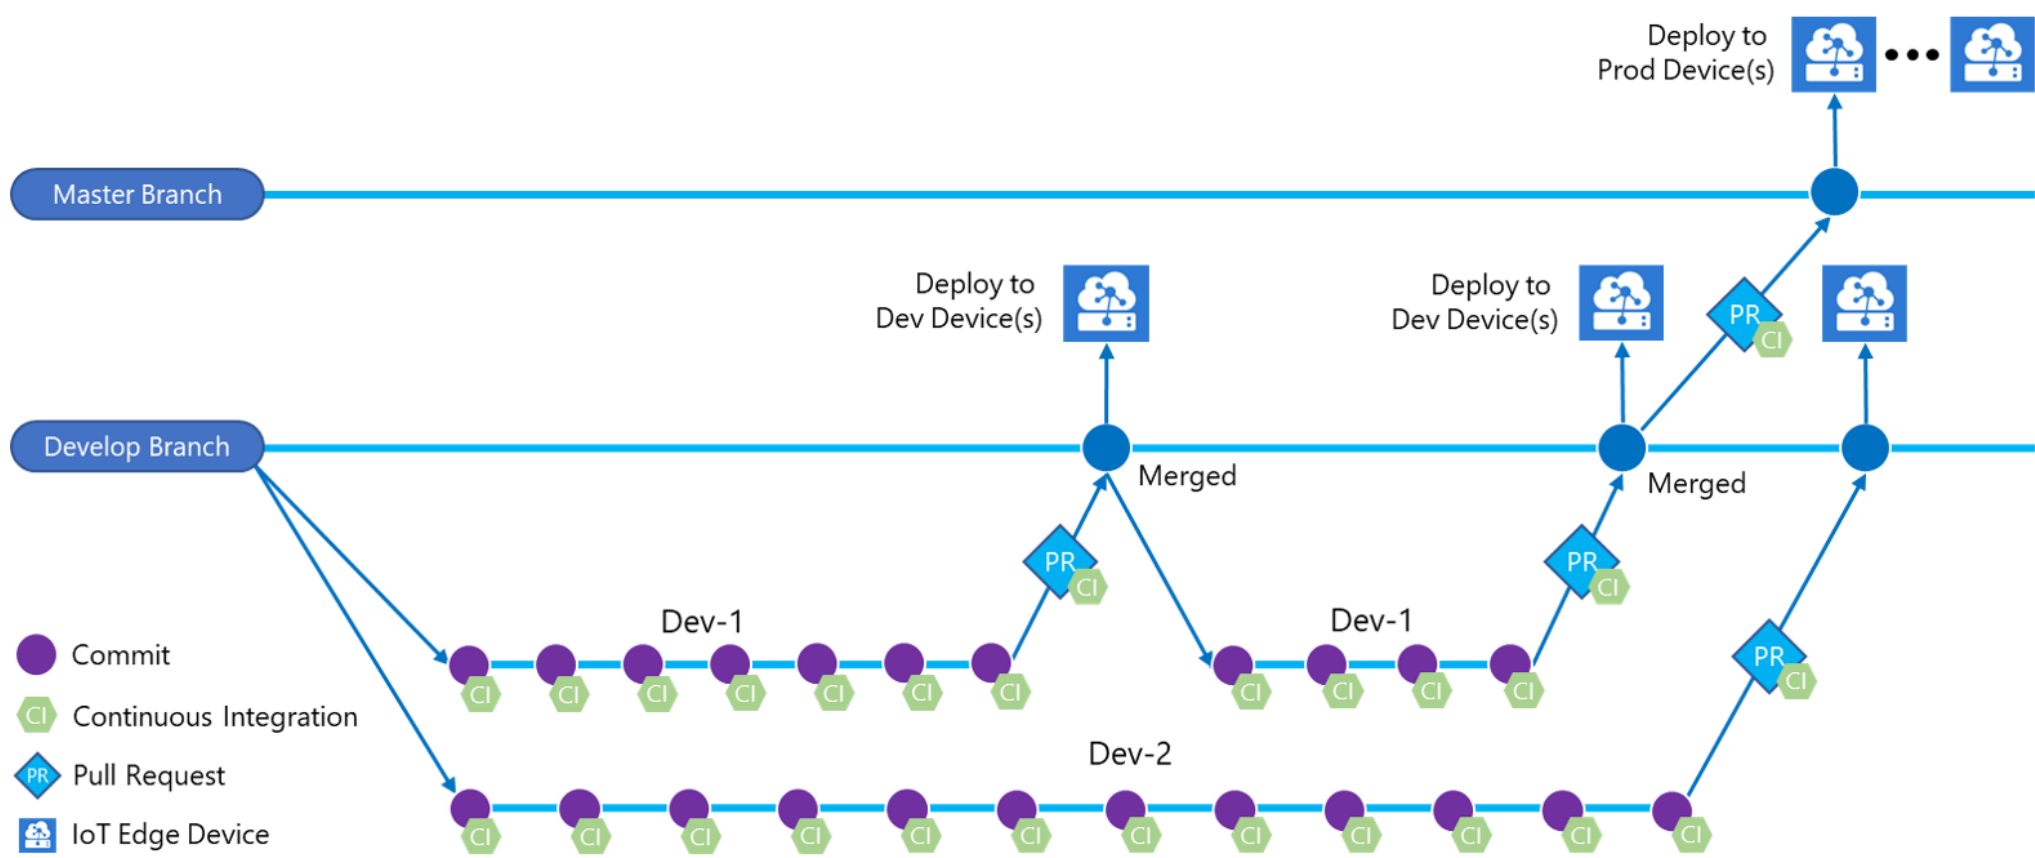
Targeted Applied Reporting Success

Filter devices

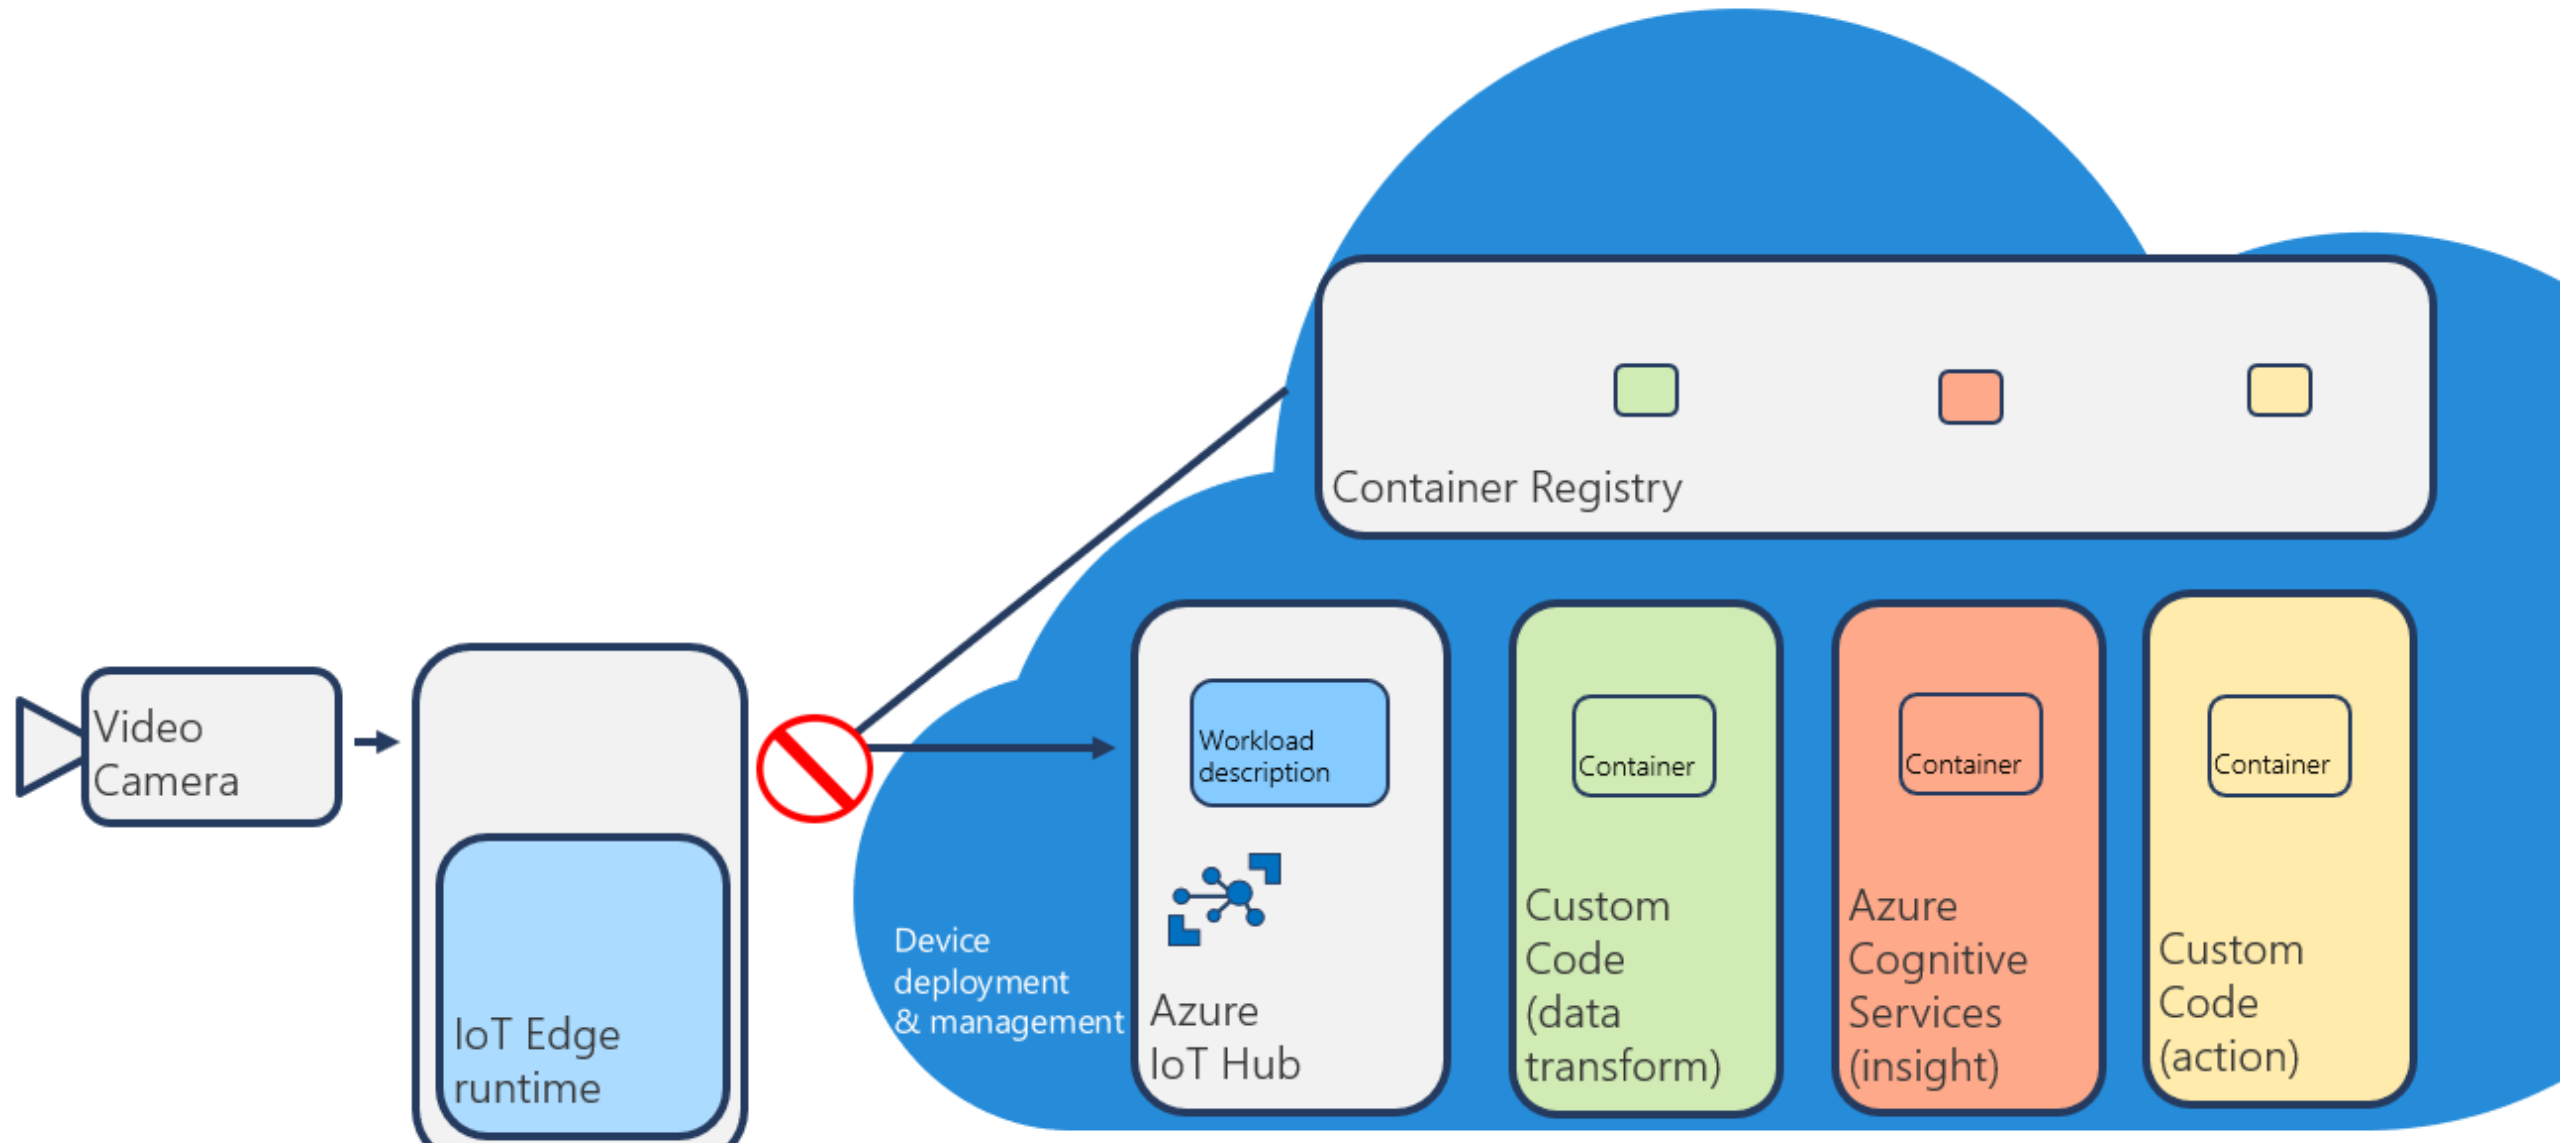
DEVICE ID
cse-paul-device
erisch

Total Devices: 2

DevOps Process with IoT Edge



Edge intelligence enabled with Azure IoT Edge



THE JETSON FAMILY

From AI at the Edge to Autonomous Machines



JETSON NANO

5–10W
0.5 TFLOPS (FP16)
45mm x 70mm



JETSON TX1 JETSON TX2 4GB

7–15W
1–1.3 TFLOPS (FP16)
50mm x 87mm



JETSON TX2 8GB | Industrial

7–15W
1.3 TFLOPS (FP16)
50mm x 87mm



JETSON AGX XAVIER

10–30W
11 TFLOPS (FP16) | 32 TOPS (INT8)
100mm x 87mm

AI at the Edge

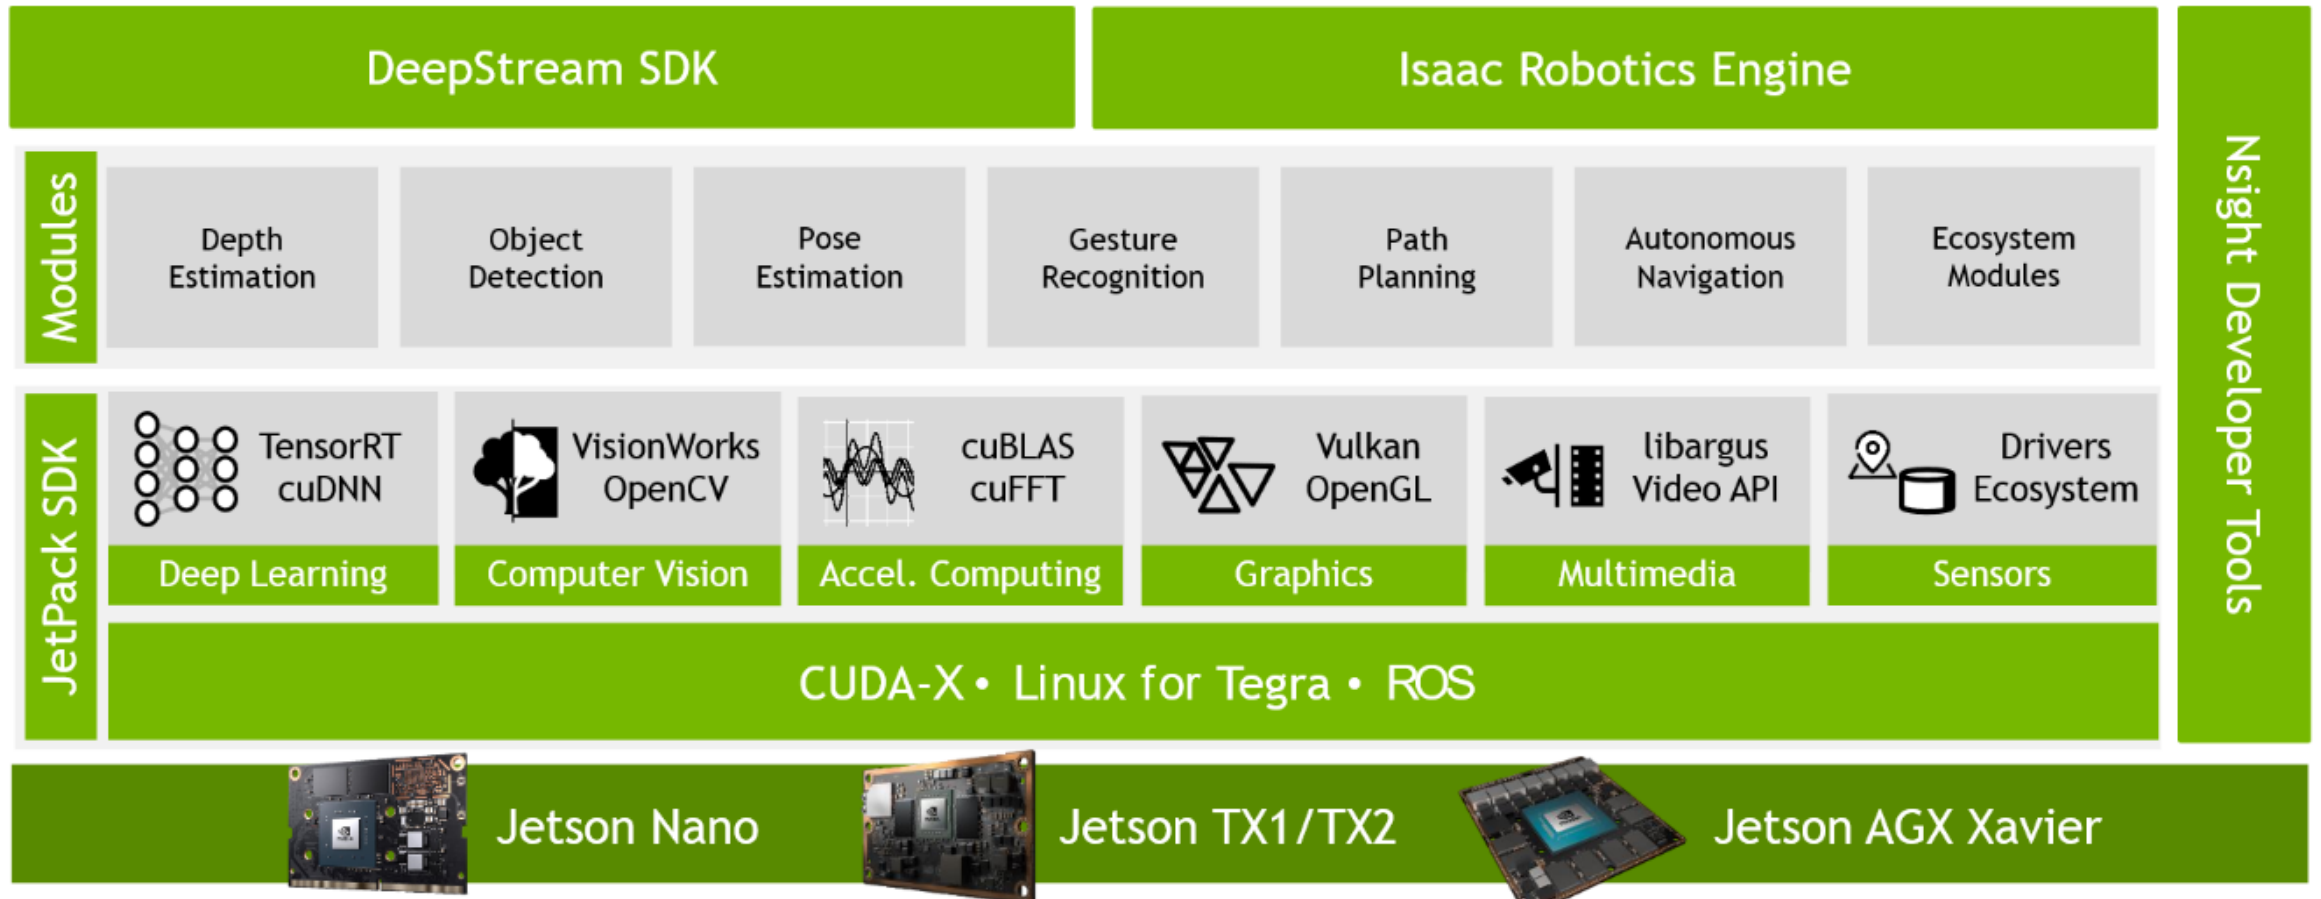
Fully Autonomous Machines

Multiple Devices – Same Software

Compare NVIDIA Jetson Module Specifications

	Jetson Nano	Jetson TX2 Series			Jetson AGX Xavier Series	
		TX2 4GB	TX2	TX2i	Jetson AGX Xavier 8GB	Jetson AGX Xavier
GPU	NVIDIA Maxwell™ architecture with 128 NVIDIA CUDA® cores	NVIDIA Pascal™ architecture with 256 NVIDIA CUDA cores			NVIDIA Volta™ architecture with 384 NVIDIA CUDA cores and 48 Tensor cores	NVIDIA Volta™ architecture with 512 NVIDIA CUDA cores and 64 Tensor cores
	0.5 TFLOPs (FP16)	1.3 TFLOPs (FP16)			5.5 TFLOPs (FP16) 11.1 TOPS (INT8)	11 TFLOPs (FP16) 22 TOPS (INT8)
CPU	Quad-core ARM® Cortex®-A57 MPCore processor	Dual-core Denver 2 64-bit CPU and quad-core ARM A57 complex			6-core Carmel ARM v8.2 64-bit CPU, 8MB L2 + 4MB L3	8-core Carmel ARM v8.2 64-bit CPU, 8MB L2 + 4MB L3

JETSON SOFTWARE



JETPACK 4.2



Available Now For Jetson
developer.nvidia.com/jetpack

Package Versions

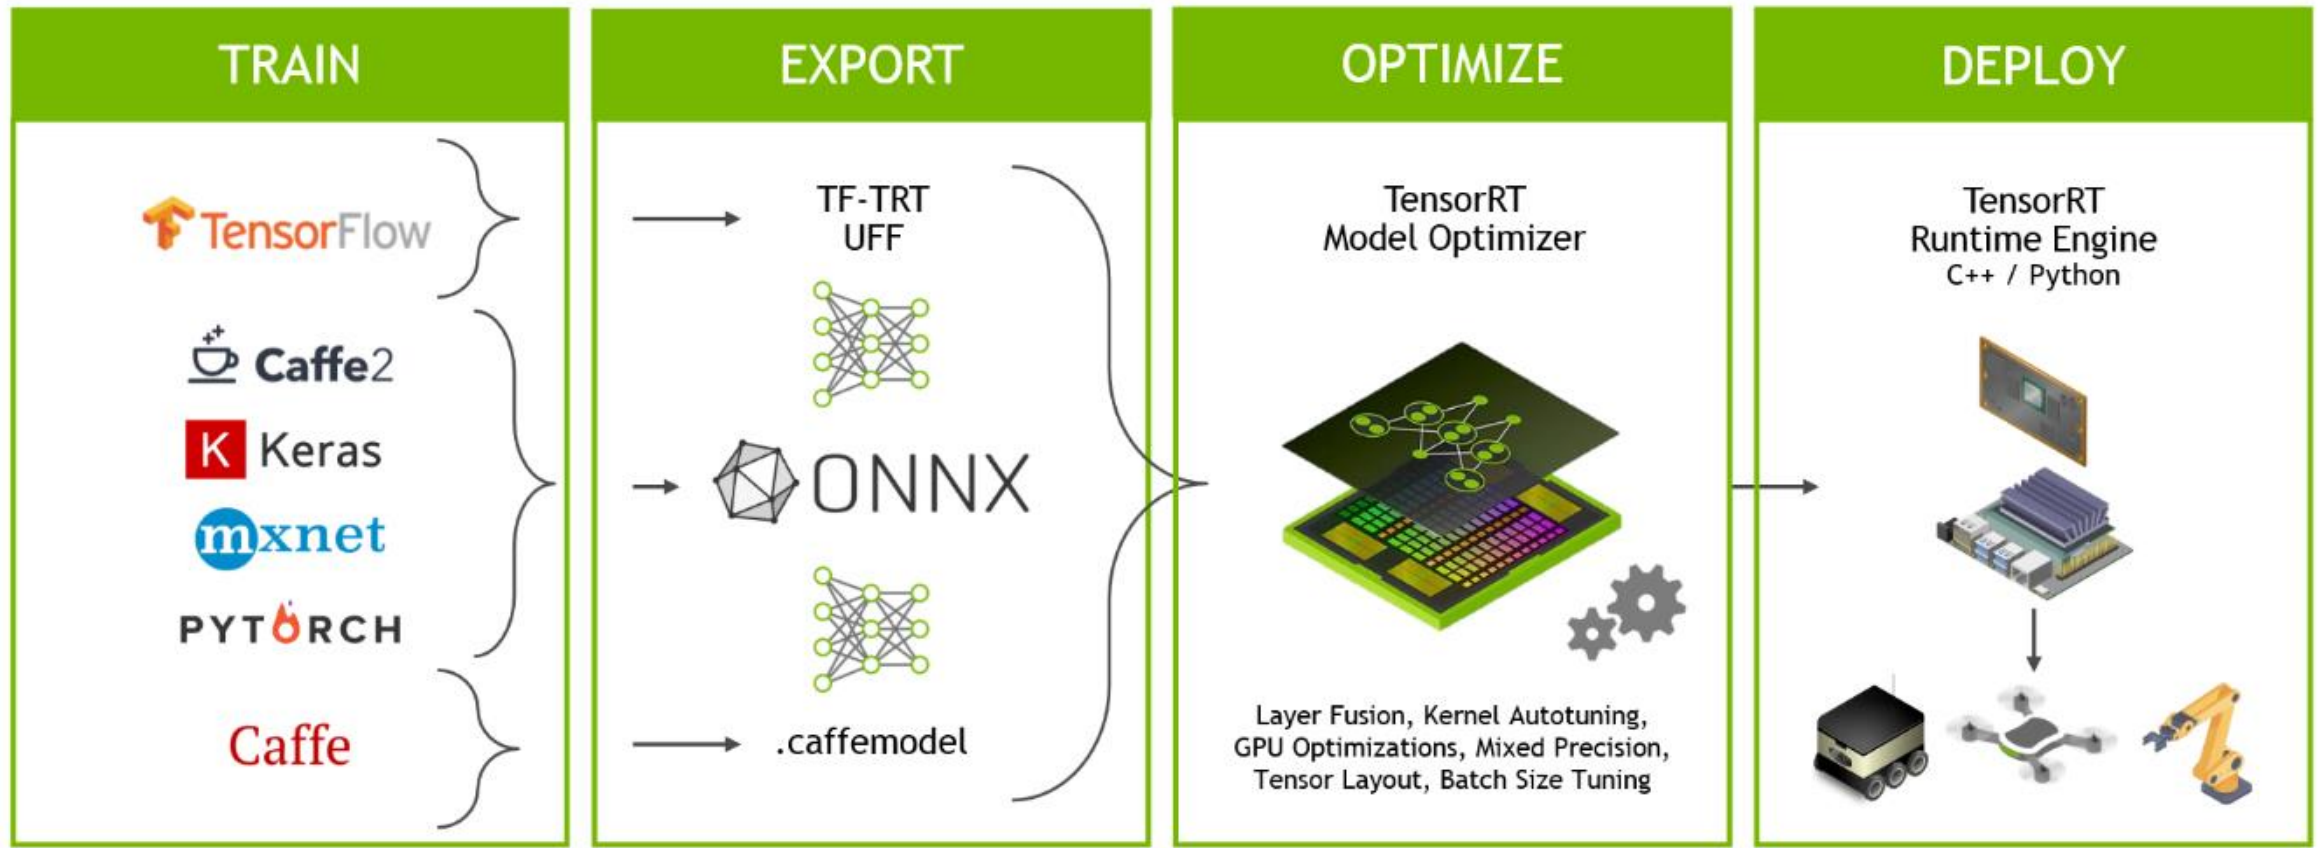
L4T BSP	32.1
Linux Kernel	4.9.140
Vulkan	1.1.1
OpenGL	4.6
OpenGL-ES	3.2.5
EGL	1.5
GLX	1.4
X11 ABI	24
Wayland	1.14
L4T Multimedia API	32.1
Argus Camera API	0.97
GStreamer	1.14.1
Nsight Systems	2019.3
Nsight Graphics	2018.7
Nsight Compute	1.0
Jetson GPIO	1.0
Jetson OS	Ubuntu 18.04
Host OS	Ubuntu 16.04 / 18.04

CUDA	10.0.166
cuDNN	7.3.1.28
TensorRT	5.0.6.3
VisionWorks	1.6
OpenCV	3.3.1
NPP	10.0


Install TensorFlow, PyTorch, Caffe,
Caffe2, MXNet, ROS, and other
GPU-accelerated libraries



NVIDIA TensorRT



JETSON NANO RUNS MODERN AI



developer.nvidia.com/embedded/jetson-nano-dl-inference-benchmarks

Achieving Real-Time Performance

- Number of 1080p/30 FPS streams processed on various platforms.


NVIDIA Products	H.264	H.265
Jetson Nano	8	8
Jetson TX1	8	8
Jetson TX2	14	14
Jetson AGX Xavier	32	49
T4	35	68

Data measured using deepstream-app from DeepStream SDK 4.0

JETSON NANO COMPUTE MODULE

PROCESSOR	
CPU	64-bit Quad-core ARM A57 @ 1.43GHz
GPU	128-core NVIDIA Maxwell @ 921MHz
Memory	4GB 64-bit LPDDR4 @ 1600MHz 25.6GB/s
Video Encoder	4Kp30 (4x) 1080p30 (2x) 1080p60
Video Decoder	4Kp60 (2x) 4Kp30 (8x) 1080p30 (4x) 1080p60


INTERFACES	
USB	USB 3.0 (3x) USB 2.0
Camera	12 lanes MIPI CSI-2 (up to 4 cameras)
Display	HDMI DP eDP DSI
Networking	Gigabit Ethernet
PCIe	PCIe Gen2 x1/x2/x4
Storage	16GB eMMC 5.1
Other I/O	(4x) I2C (2x) SPI (3x) UART (2x) I2S GPIO
Power	5V DC, 5W 10W
Size	45x70mm, 260-pin SODIMM connector



Production module
available June 2019

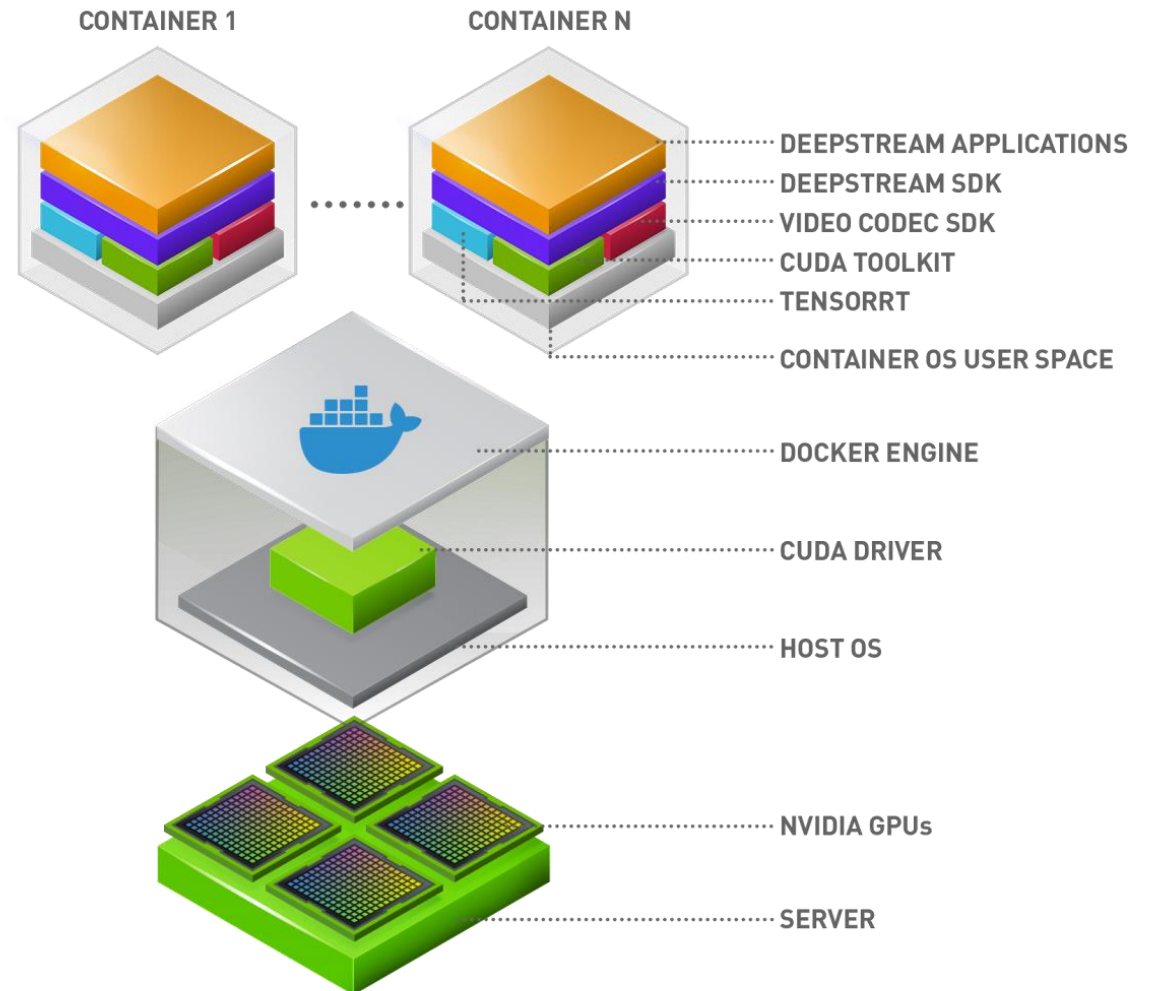
Nvidia Docker

- CUDA Driver resides as part of the Host OS
- At runtime, nvidia-docker mounts CUDA Driver and adds devices to container
- Allows container to leverage GPU capabilities



Deepstream Module

- Works on same principal as Nvidia Docker
- Contains Deepstream Application which is driven by a config file
- Supports messaging to Azure IoT Hub from Deepstream App



Demo

Nvidia Deepstream on Jetson Xavier

Video

Vision IQ Solution by the Marsden Group

Further Learning



Additional Resources

Onnx-runtime-iotedge:

<https://github.com/Azure-Samples/onnxruntime-iot-edge>

- Use onnx models with webcam input on Nano with IoT Edge

Ian Davis's blog:

<https://codepyre.com/>

- Contains info on building custom RootFS / Tensorflow

Paul DeCarlo's Dev.to articles:

<https://dev.to/toolboc>

- Articles related to IoT Edge Development on Nvidia devices

This powerpoint presentation:

<http://aka.ms/IntelligentEdgeHOLdeck>

Video recording of this presentation:

<http://aka.ms/IntelligentEdgeHOLvideo>

Hardware Checklist



Bill of Materials

Jetson Nano DevKit - <https://amzn.to/2WFE5zF>

Jetson Nano Case + Fan - <https://amzn.to/2ZI2ki9>

Jetson Nano 4V 2A barrel adapter - <https://amzn.to/32DFsTq>

Micro Jumper - <https://amzn.to/2SoyPQa>

128GB MicroSD card - <https://amzn.to/32zj0ep>

Edimax WiFi Adapter - <https://amzn.to/2SuDXT8>

